Zhisheng_Lin_HW1

January 31, 2020

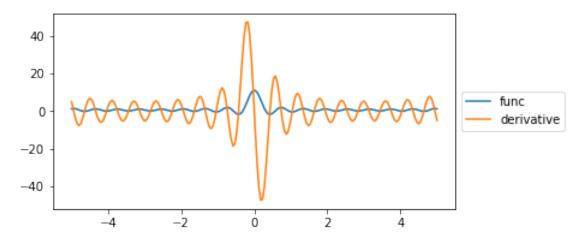
Exercise 1. Python class refresher

```
[53]: import autograd.numpy as np
      import matplotlib.pyplot as plt
      from autograd import grad
      class GradViewer():
          def __init__(self,func):
              self.func = func
          def plot_it(self):
              # create space over which to evaluate function and gradient
              w_vals = np.linspace(-5,5,200)
              # get function
              g = self.func
              # calculate the gradient function
              nabla_g = grad(g)
              # evaluate gradient over input range
              g_vals = [g(v) for v in w_vals]
              grad_vals = [nabla_g(v) for v in w_vals]
              # create figure
              fig, ax = plt.subplots(1, 1, figsize=(6,3))
              # plot function and gradient values
              ax.plot(w_vals,g_vals)
              ax.plot(w_vals,grad_vals)
              ax.legend(['func','derivative'],loc='center left', bbox_to_anchor=(1, 0.
       →5))
              plt.show()
      # test function input
      def my_function(w):
          # initialize y for storing the sum
          y = 0
```

```
# use for loop to realize summing
for i in range(11):
    y = y+np.cos(i*w)
    return y

# create instance, inputting function on creation
test = GradViewer(my_function)

# plot this derivative and original function
test.plot_it()
```



Exercise 2. Python class refresher part 2

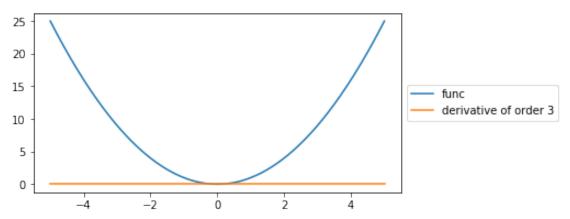
```
[66]: import autograd.numpy as np
  import matplotlib.pyplot as plt
  from autograd import grad

class OrderViewer(GradViewer):
    def __init__(self,func):
        super().__init__(func)

    def compute_it(self,order):
        global deri
        deri = self.func
        for i in range(int(order)):
            deri = grad(deri)
        return deri

    def plot_it(self):
        # create space over which to evaluate function and gradient
        w_vals = np.linspace(-5,5,200)
```

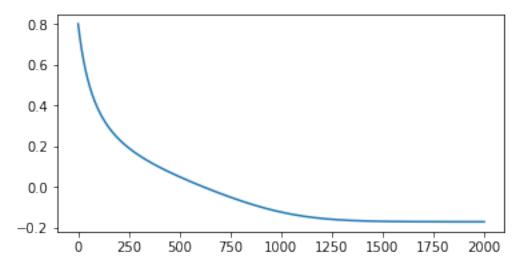
```
# get function
        g = self.func
        # calculate the gradient function
        nabla_g = deri
        # evaluate gradient over input range
        g_{vals} = [g(v) \text{ for } v \text{ in } w_{vals}]
        grad_vals = [nabla_g(v) for v in w_vals]
        # create figure
        fig, ax = plt.subplots(1, 1, figsize=(6,3))
        # plot function and gradient values
        ax.plot(w_vals,g_vals)
        ax.plot(w_vals,grad_vals)
        ax.legend(['func','derivative of order 3'],loc='center left',u
 \rightarrowbbox_to_anchor=(1, 0.5))
        plt.show()
# test function input
def my_function(w):
    y = w**2
    return y
# create instance, inputting function on creation
test = OrderViewer(my_function)
# compute desired order derivative
test.compute_it(order = 3)
# plot this derivative and original function
test.plot_it()
```



Exercise 3. A generic gradient descent function Input Function 1: $g = \frac{1}{50}(w^4 + w^2 + 10w)$

```
[115]: import autograd.numpy as np
       import matplotlib.pyplot as plt
       from autograd import grad
       # gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
        →parameter), max_its (maximum number of iterations), w (initialization)
       def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
           gradient = grad(g)
           # run the gradient descent loop
           best_w = w # weight we return, should be the one providing lowest_
        \rightarrow evaluation
           best_eval = g(w)  # lowest evaluation yet
best_w_arr = []  # store weights in each step
           best_w_arr.append(best_w) # store the first weight
           for k in range(max_its):
               # evaluate the gradient
               grad_eval = gradient(w)
               # take gradient descent step
               w = w - alpha*grad_eval
               # return only the weight providing the lowest evaluation
               test_eval = g(w)
               if test_eval < best_eval:</pre>
                    best_eval = test_eval
                    best_w = w
               best_w_arr.append(best_w)
           return best_w_arr
       # create the input function
       g = lambda w: 1/float(50)*(w**4 + w**2 + 10*w) # try other functions too! <math>\Box
        \rightarrowLike g = lambda w: np.cos(2*w) , g = lambda w: np.sin(5*w) + 0.1*w**2, g =\Box
        \rightarrow lambda w: np.cos(5*w)*np.sin(w)
       # run gradient descent
       weight_history = gradient_descent(g = g,alpha = 10**-2,max_its = 2000,w = 2.0)
       # cost function history plotter
       def cost_history(weight_history,g):
```

```
# loop over weight history and compute associated cost function history at \Box
 →each step
    weight_history = np.asarray(weight_history)
    cost_history = g(weight_history)
    # plot cost function history
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,2000,len(cost_history)),cost_history)
    plt.show()
    # compare the weight corresponding to the smallest cost value with the \Box
 \rightarrow expected weight
    small_costval_arg = np.argmin(cost_history)
    small_weightval = weight_history[small_costval_arg]
    w_{exp} = ((2031**(1/2)-45)**(1/3))/(6**(2/3))-1/(((2031**(1/2)-45)*6)**(1/3))
    print("The weight corresponding to the smallest cost function value is ⊔
 \rightarrow", end='')
    print(small_weightval)
    print("The expected weight is ",end='')
    print(w_exp)
    print("So, they are close.")
# use cost_history plotter
cost_history(weight_history,g)
```



The weight corresponding to the smallest cost function value is -1.2198339861245202

The expected weight is -1.2347728250533112

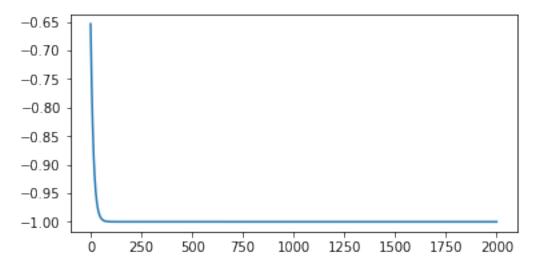
So, they are close. $Input \ Function \ 2: \ g = cos(2w)$

```
[112]: import autograd.numpy as np
       import matplotlib.pyplot as plt
       from autograd import grad
       # gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
        →parameter), max_its (maximum number of iterations), w (initialization)
       def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
           gradient = grad(g)
           # run the gradient descent loop
           best_w = w
                         # weight we return, should be the one providing lowest
        \rightarrow evaluation
           best_eval = g(w)  # lowest evaluation yet
best_w_arr = []  # store weights in each step
                                             # store the first weight
           best_w_arr.append(best_w)
           for k in range(max_its):
                # evaluate the gradient
                grad_eval = gradient(w)
                # take gradient descent step
                w = w - alpha*grad_eval
                # return only the weight providing the lowest evaluation
                test_eval = g(w)
                if test_eval < best_eval:</pre>
                    best_eval = test_eval
                    best_w = w
                best_w_arr.append(best_w)
           return best w arr
       # create the input function
       g = lambda w: np.cos(2*w) # try other functions too! Like <math>g = lambda w: np.
        \rightarrow \cos(2*w), q = lambda w: np.sin(5*w) + 0.1*w**2, <math>q = lambda w: np.cos(5*w)*np.
        \rightarrow sin(w)
       # run gradient descent
       weight_history = gradient_descent(g = g,alpha = 10**-2,max_its = 2000,w = 2.0)
       # cost function history plotter
       def cost_history(weight_history,g):
           # loop over weight history and compute associated cost function history at 11
        \rightarrow each step
```

```
weight_history = np.asarray(weight_history)
cost_history = g(weight_history)

# plot cost function history
# create figure
fig, ax = plt.subplots(1, 1, figsize=(6,3))
# plot function and gradient values
ax.plot(np.linspace(0,2000,len(cost_history)),cost_history)
plt.show()

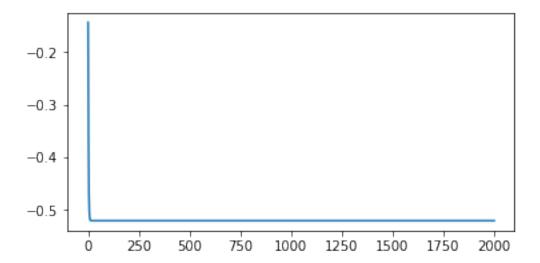
# use cost_history plotter
cost_history(weight_history,g)
```



Input Function 3: $g = \sin(5w) + 0.1w^2$

```
[113]: import autograd.numpy as np
       import matplotlib.pyplot as plt
       from autograd import grad
       \# gradient descent function - inputs: g (input function), alpha (steplength_\sqcup
        →parameter), max_its (maximum number of iterations), w (initialization)
       def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
           gradient = grad(g)
           # run the gradient descent loop
           best_w = w
                            # weight we return, should be the one providing lowest
        \rightarrow evaluation
           best_eval = g(w)
                                  # lowest evaluation yet
           best_w_arr = []
                                  # store weights in each step
```

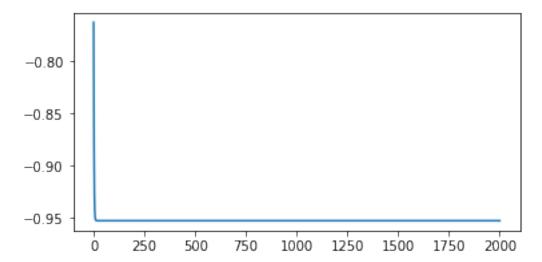
```
best_w_arr.append(best_w)
                                     # store the first weight
    for k in range(max_its):
        # evaluate the gradient
        grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
        if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
        best_w_arr.append(best_w)
    return best_w_arr
# create the input function
g = lambda w: np.sin(5*w) + 0.1*w**2 # try other functions too! Like <math>g = lambda
\rightarrowlambda w: np.cos(2*w) , g = lambda w: np.sin(5*w) + 0.1*w**2, g = lambda w: np.
\rightarrow \cos(5*w)*np.\sin(w)
# run gradient descent
weight_history = gradient_descent(g = g,alpha = 10**-2,max_its = 2000,w = 2.0)
# cost function history plotter
def cost_history(weight_history,g):
    \# loop over weight history and compute associated cost function history at \sqcup
→each step
    weight_history = np.asarray(weight_history)
    cost_history = g(weight_history)
    # plot cost function history
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,2000,len(cost_history)),cost_history)
    plt.show()
# use cost_history plotter
cost_history(weight_history,g)
```



Input Function 4: g = cos(5w)sin(w)

```
[146]: import autograd.numpy as np
       import matplotlib.pyplot as plt
       from autograd import grad
       # gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
        →parameter), max_its (maximum number of iterations), w (initialization)
       def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
           gradient = grad(g)
           # run the gradient descent loop
                              # weight we return, should be the one providing lowest \square
           best_w = w
        \rightarrow evaluation
                                   # lowest evaluation yet
           best_eval = g(w)
           best_w_arr = []
                                  # store weights in each step
           best_w_arr.append(best_w)
                                            # store the first weight
           for k in range(max_its):
               # evaluate the gradient
               grad_eval = gradient(w)
               # take gradient descent step
               w = w - alpha*grad_eval
               # return only the weight providing the lowest evaluation
               test_eval = g(w)
               if test_eval < best_eval:</pre>
                   best_eval = test_eval
                   best_w = w
```

```
best_w_arr.append(best_w)
    return best_w_arr
# create the input function
g = lambda w: np.cos(5*w)*np.sin(w) # try other functions too! Like <math>g = lambda
 \rightarrowlambda w: np.cos(2*w) , q = lambda w: np.sin(5*w) + 0.1*w**2, q = lambda w: np.
\hookrightarrow \cos(5*w)*np.\sin(w)
# run gradient descent
weight_history = gradient_descent(g = g,alpha = 10**-2,max_its = 2000,w = 2.0)
# cost function history plotter
def cost_history(weight_history,g):
    # loop over weight history and compute associated cost function history at \square
 \rightarrow each step
    weight_history = np.asarray(weight_history)
    cost_history = g(weight_history)
    # plot cost function history
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,2000,len(cost_history)),cost_history)
    plt.show()
# use cost_history plotter
cost_history(weight_history,g)
```

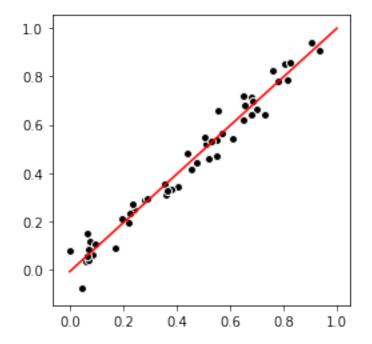


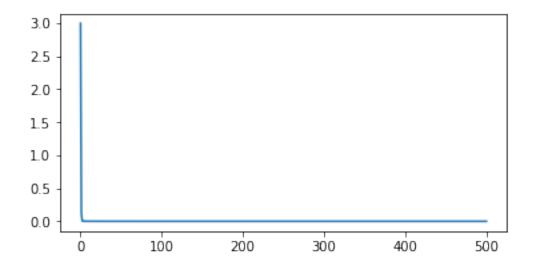
Exercise 4. Apply gradient descent to minimize the Least Squares cost for linear regression on a low dimensional dataset $1. \alpha_1 = 0.01$

```
[161]: import autograd.numpy as np
      import matplotlib.pyplot as plt
      from autograd import grad
       # data input
      csvname = '2d_linregress_data.csv'
      data = np.loadtxt(csvname,delimiter = ',')
      # form the input/output data vectors
      x = data[:,:-1]
      y = data[:,-1]
       # least squares cost function for linear regression
      def least_squares(w):
          cost = 0
           for p in range(len(y)):
              # get pth input/output pair
              x_p = x[p]
              y_p = y[p]
               # form linear combination
               c_p = w[0] + w[1]*x_p
               # add least squares for this datapoint
               cost += (c_p - y_p)**2
          return cost
       # gradient descent function - inputs: g (input function), alpha (steplengthu
        →parameter), max_its (maximum number of iterations), w (initialization)
      def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
           gradient = grad(g)
           # run the gradient descent loop
           best_w = w # weight we return, should be the one providing lowest_\(\sigma\)
        \rightarrow evaluation
          best_eval = g(w)
                                # lowest evaluation yet
          best_w_arr = [] # store weights in each step
          best_w_arr.append(best_w) # store the first weight
           for k in range(max_its):
               # evaluate the gradient
```

```
grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
        if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
        best_w_arr.append(best_w)
    return best_w_arr
# run gradient descent
w = np.asarray([1.5, 1.5])
weight_history = gradient_descent(g = least_squares,alpha = 10**-2,max_its = __
\rightarrow 500, w = w)
weight_history = np.asarray(weight_history)
# MSE history plotter
def MSE(weight_history,g):
    MSE_arr = []
    # loop over weight history and compute the MSE at each step o gradient_{\square}
 \rightarrow descent
    for i in range(len(weight_history)):
        MSE = g(weight_history[i])/len(y)
        MSE_arr.append(MSE)
    # plot MSE
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,500,len(MSE_arr)),MSE_arr)
    plt.show()
# scatter plot the input data
fig, ax = plt.subplots(1, 1, figsize=(4,4))
ax.scatter(data[:,0],data[:,1],color = 'k',edgecolor = 'w')
# fit a trend line
x_vals = np.linspace(0,1,200)
y_vals = weight_history[-1,0] + weight_history[-1,1]*x_vals
ax.plot(x_vals,y_vals,color = 'r')
plt.show()
# plot MSE
```

MSE(weight_history,least_squares)



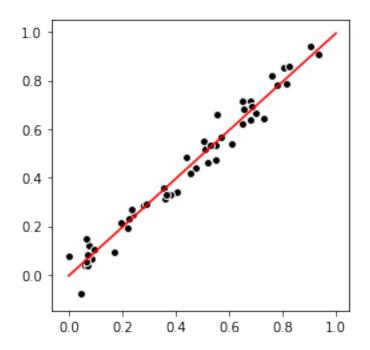


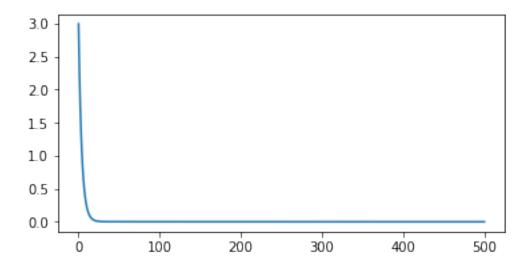
2. $\alpha_2 = 0.001$

[2]: import autograd.numpy as np import matplotlib.pyplot as plt from autograd import grad

```
# data input
csvname = '2d_linregress_data.csv'
data = np.loadtxt(csvname,delimiter = ',')
# form the input/output data vectors
x = data[:,:-1]
y = data[:,-1]
# least squares cost function for linear regression
def least_squares(w):
   cost = 0
    for p in range(len(y)):
        # get pth input/output pair
       x_p = x[p]
       y_p = y[p]
        # form linear combination
        c_p = w[0] + w[1]*x_p
        # add least squares for this datapoint
        cost += (c_p - y_p)**2
    return cost
# gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
→parameter), max_its (maximum number of iterations), w (initialization)
def gradient_descent(g,alpha,max_its,w):
    # compute the gradient of our input function - note this is a function too!
    gradient = grad(g)
    # run the gradient descent loop
   best_w = w # weight we return, should be the one providing lowest_
 \rightarrow evaluation
    best_eval = g(w)
                           # lowest evaluation yet
    best_w_arr = [] # store weights in each step
                                # store the first weight
    best_w_arr.append(best_w)
    for k in range(max_its):
        # evaluate the gradient
        grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
```

```
if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
        best_w_arr.append(best_w)
    return best_w_arr
# run gradient descent
w = np.asarray([1.5, 1.5])
weight_history = gradient_descent(g = least_squares,alpha = 10**-3,max_its =__
\rightarrow500, w = w)
weight_history = np.asarray(weight_history)
# MSE history plotter
def MSE(weight_history,g):
    MSE_arr = []
    \# loop over weight history and compute the MSE at each step o gradient_{\sqcup}
\rightarrow descent
    for i in range(len(weight_history)):
        MSE = g(weight_history[i])/len(y)
        MSE_arr.append(MSE)
    # plot MSE
    # create figure
   fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,500,len(MSE_arr)),MSE_arr)
    plt.show()
# scatter plot the input data
fig, ax = plt.subplots(1, 1, figsize=(4,4))
ax.scatter(data[:,0],data[:,1],color = 'k',edgecolor = 'w')
# fit a trend line
x_vals = np.linspace(0,1,200)
y_vals = weight_history[-1,0] + weight_history[-1,1]*x_vals
ax.plot(x_vals,y_vals,color = 'r')
plt.show()
# plot MSE
MSE(weight_history,least_squares)
```





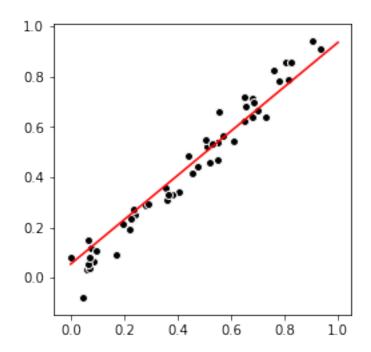
```
3. \alpha_3 = 0.0001
```

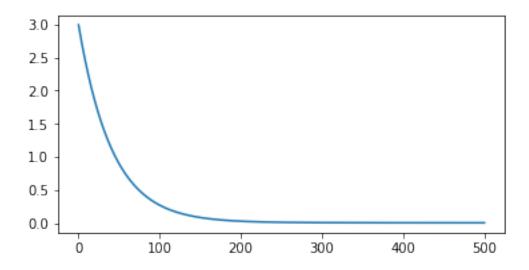
```
[3]: import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# data input
csvname = '2d_linregress_data.csv'
data = np.loadtxt(csvname,delimiter = ',')
```

```
# form the input/output data vectors
x = data[:,:-1]
y = data[:,-1]
# least squares cost function for linear regression
def least_squares(w):
    cost = 0
    for p in range(len(y)):
        # get pth input/output pair
       x_p = x[p]
       y_p = y[p]
        # form linear combination
        c_p = w[0] + w[1]*x_p
        # add least squares for this datapoint
        cost += (c_p - y_p)**2
   return cost
# gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
→parameter), max_its (maximum number of iterations), w (initialization)
def gradient_descent(g,alpha,max_its,w):
   # compute the gradient of our input function - note this is a function too!
    gradient = grad(g)
    # run the gradient descent loop
                 # weight we return, should be the one providing lowest
    best_w = w
 \rightarrow evaluation
   best_eval = g(w)
                         # lowest evaluation yet
   best_w_arr = [] # store weights in each step
    best_w_arr.append(best_w)
                                 # store the first weight
    for k in range(max_its):
        # evaluate the gradient
        grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
        if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
```

```
best_w_arr.append(best_w)
    return best_w_arr
# run gradient descent
w = np.asarray([1.5, 1.5])
weight_history = gradient_descent(g = least_squares,alpha = 10**-4,max_its =__
\rightarrow500, w = w)
weight_history = np.asarray(weight_history)
# MSE history plotter
def MSE(weight_history,g):
   MSE_arr = []
    # loop over weight history and compute the MSE at each step o gradient_{\square}
 \rightarrowdescent
    for i in range(len(weight_history)):
        MSE = g(weight_history[i])/len(y)
        MSE_arr.append(MSE)
    # plot MSE
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,500,len(MSE_arr)),MSE_arr)
    plt.show()
# scatter plot the input data
fig, ax = plt.subplots(1, 1, figsize=(4,4))
ax.scatter(data[:,0],data[:,1],color = 'k',edgecolor = 'w')
# fit a trend line
x_vals = np.linspace(0,1,200)
y_vals = weight_history[-1,0] + weight_history[-1,1]*x_vals
ax.plot(x_vals,y_vals,color = 'r')
plt.show()
# plot MSE
MSE(weight_history,least_squares)
```





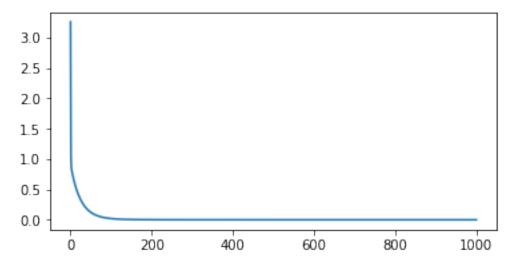
Exercise 5. Use linear regression to fit to a high dimensional dataset $1. \alpha = 0.001$, iterations = 1000

```
[2]: import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# load in dataset
```

```
data = np.loadtxt('highdim_linregress_data.csv',delimiter = ',')
# form the input/output data vectors
x = data[:,:-1]
y = data[:,-1]
# change y to a array
y = np.asarray(y).reshape(100,1)
\# concatenate 11 by 1 all one vector on the left of input x
o = np.ones((np.shape(x)[0],1))
x = np.concatenate((o,x),axis=1)
# least squares cost function for linear regression
def least_squares(w):
   # get the calculation value
   c_p = x.dot(w)
    # get the real value
   y_p = y
   # get the cost
    cost = np.sum((c_p-y_p)**2)
   return cost
# gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
→parameter), max_its (maximum number of iterations), w (initialization)
def gradient_descent(g,alpha,max_its,w):
    # compute the gradient of our input function - note this is a function too!
    gradient = grad(g)
    # run the gradient descent loop
   best_w = w # weight we return, should be the one providing lowest_
 \rightarrow evaluation
    best_eval = g(w)
                         # lowest evaluation yet
    best_w_arr = [] # store weights in each step
                                # store the first weight
    best_w_arr.append(best_w)
    for k in range(max_its):
        # evaluate the gradient
        grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
```

```
if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
        best_w_arr.append(best_w)
    return best_w_arr
# run gradient descent
w = np.random.randn(np.shape(x)[1],1)
weight_history = gradient_descent(g = least_squares,alpha = 10**-3,max_its =
\rightarrow 1000, w = w
weight_history = np.asarray(weight_history)
# MSE history plotter
def MSE(weight_history,g):
    MSE_arr = []
    # loop over weight history and compute the MSE at each step o gradient
 \rightarrow descent
    for i in range(len(weight_history)):
        MSE = g(weight_history[i])/len(y)
        MSE_arr.append(MSE)
    # plot MSE
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,1000,len(MSE_arr)),MSE_arr)
    plt.show()
# plot MSE
MSE(weight_history,least_squares)
```



```
2. \alpha = 0.0001, iterations = 100
```

```
[189]: import autograd.numpy as np
      import matplotlib.pyplot as plt
      from autograd import grad
      # load in dataset
      data = np.loadtxt('highdim_linregress_data.csv',delimiter = ',')
       # form the input/output data vectors
      x = data[:,:-1]
      y = data[:,-1]
      # change y to a array
      y = np.asarray(y).reshape(100,1)
      \# concatenate 11 by 1 all one vector on the left of input x
      o = np.ones((np.shape(x)[0],1))
      x = np.concatenate((o,x),axis=1)
       # least squares cost function for linear regression
      def least_squares(w):
          # get the calculation value
          c_p = x.dot(w)
          # get the real value
          y_p = y
          # get the cost
          cost = np.sum((c_p-y_p)**2)
          return cost
       # gradient descent function - inputs: g (input function), alpha (steplength_{\sqcup}
       →parameter), max_its (maximum number of iterations), w (initialization)
      def gradient_descent(g,alpha,max_its,w):
           # compute the gradient of our input function - note this is a function too!
          gradient = grad(g)
          # run the gradient descent loop
          best_w = w # weight we return, should be the one providing lowest_u
        \rightarrow evaluation
          best_eval = g(w)
                                # lowest evaluation yet
          best_w_arr = [] # store weights in each step
          best_w_arr.append(best_w) # store the first weight
```

```
for k in range(max_its):
        # evaluate the gradient
        grad_eval = gradient(w)
        # take gradient descent step
        w = w - alpha*grad_eval
        # return only the weight providing the lowest evaluation
        test_eval = g(w)
        if test_eval < best_eval:</pre>
            best_eval = test_eval
            best_w = w
        best_w_arr.append(best_w)
    return best_w_arr
# run gradient descent
w = np.random.randn(np.shape(x)[1],1)
weight_history = gradient_descent(g = least_squares,alpha = 10**-4,max_its = 10**-4.
\rightarrow 100, w = w
weight_history = np.asarray(weight_history)
# MSE history plotter
def MSE(weight_history,g):
   MSE_arr = []
    # loop over weight history and compute the MSE at each step o gradient_{\sqcup}
 \rightarrow descent
    for i in range(len(weight_history)):
       MSE = g(weight_history[i])/len(y)
        MSE_arr.append(MSE)
    # plot MSE
    # create figure
    fig, ax = plt.subplots(1, 1, figsize=(6,3))
    # plot function and gradient values
    ax.plot(np.linspace(0,100,len(MSE_arr)),MSE_arr)
    plt.show()
# plot MSE
MSE(weight_history,least_squares)
```

